

Training Loss Curves — $\beta=50$

(dense spikes = collocation resampling [Pattern A]; large isolated spikes = FF gradient instability [Pattern B])

