

Loss Sharpness Comparison (mean $\pm 2\sigma$, n=5 runs)

Counter-intuitive result: failed model has SHARPER landscape

Hessian Max Eigenvalue $\lambda_{\max}$

