

Training Loss Trajectories by Dimensionality ($\beta=30$)

Architecture: 4×64 tanh, Adam lr=0.001

